TAMPEREEN TEKNILLINEN YLIOPISTO
TAMPERE UNIVERSITY OF TECHNOLOGY

JAAKKO PASANEN
NATURAL LANGUAGE SYNTACTIC PARSING WITH DEEP LEARNING

Master of Science thesis

ABSTRACT

The abstract is a concise 1-page description of the work: what was the problem, what was done, and what are the results. Do not include charts or tables in the abstract.

Put the abstract in the primary language of your thesis first and then the translation (when that is needed).

TIIVISTELMÄ

The abstract in Finnish. Foreign students do not need this page.

Suomenkieliseen diplomityöhön kirjoitetaan tiivistelmä sekä suomeksi että englanniksi.

Kandidaatintyön tiivistelmä kirjoitetaan ainoastaan kerran, samalla kielellä kuin työ. Kuitenkin myös suomenkielisillä kandidaatintöillä pitää olla englanninkielinen otsikko arkistointia varten.

PREFACE

This document template conforms to Guide to Writing a Thesis at Tampere University of Technology (2014) and is based on the previous template. The main purpose is to show how the theses are formatted using LaTeX (or LaTeX to be extra fancy) .

The thesis text is written into file `d_tyo.tex`, whereas `tutthesis.cls` contains the formatting instructions. Both files include lots of comments (start with %) that should help in using LaTeX. TUT specific formatting is done by additional settings on top of the original `report.cls` class file. This example needs few additional files: TUT logo, example figure, example code, as well as example bibliography and its formatting (`.bst`) An example makefile is provided for those preferring command line. You are encouraged to comment your work and to keep the length of lines moderate, e.g. <80 characters. In Emacs, you can use `Alt-Q` to break long lines in a paragraph and `Tab` to indent commands (e.g. inside figure and table environments). Moreover, tex files are well suited for versioning systems, such as Subversion or Git.

Acknowledgements to those who contributed to the thesis are generally presented in the preface. It is not appropriate to criticize anyone in the preface, even though the preface will not affect your grade. The preface must fit on one page. Add the date, after which you have not made any revisions to the text, at the end of the preface.

Tampere, 11.8.2014

On behalf of the working group, Erno Salminen

CONTENTS

LIST OF ABBREVIATIONS AND SYMBOLS

ANN	Artificial Neural Network
LAS	Labelled Attachment Score (% of tokens with correct dependency head and relation)
LDA	Latent Dirilecht Allocation
LSA	Latent Semantic Analysis
LSTM	Long short term memory; type of RNN with short term memory.
NER	Named entity recognition
NLP	Natural Language Processing
PMI	Pointwise mutual information
POS	Part-of-speech; also called lexical category
RNN	Recurrent neural network
S-LSTM	Stack long short term memory
TUT	Tampere University of Technology
UAS	Unlabelled Attachment Score (% of tokens with correct dependency head)

NOTES

0.1 Reading

These vectors can be used as features in a variety of applications, such as information retrieval (Manning et al., 2008), document classification (Sebastiani, 2002), question answering (Tellex et al., 2003). Pennington et al. 2014

See first paragraph of section 1 of Liang et al. 2016 for sources on machine translation

Encoder-decoder model

0.2 Terms for Computational Linguistics

1-of-V Coding Representing words as sparse binary vectors which have 1 at the word's vocabulary index and 0 all others. With vocabulary {dog, cat, mouse}, dog becomes [1, 0, 0], cat becomes [0, 1, 0] and mouse [0, 0, 1].

Bag-of-words Multiset of words appearing in a text with occurrence counts for each word. Used as a tool for feature generation. Does not preserve word order or grammar. Can implemented as a dictionary (or associative array) where words are the keys and counts are the values.

Conditional Random Field

Constituent In syntactic analysis, a constituent is a word or a group of words that function(s) as a single unit within a hierarchical structure. Many constituents are phrases. *Yesterday I saw **an orange bird with a white neck***

Corpus A collection of texts with linguistic annotations.

Dimensionality When discussing word embeddings and word vector spaces the dimensionality refers to definition in linear algebra. Dimensionality of arrays in computing means the number of indices required to specify an element in the array. Word vector in 50 dimensional vector space $\mathbb{R}^{50}$ would be represented in computing as one dimensional array of length 50 [d1, d2, d3, ..., d50]

Distributional Hypothesis Words that are used and occur in the same contexts tend to purport similar meanings. Harris 1954

Feature Numeric data representation that can be effectively exploited in machine learning tasks. E.g. Word occurrence frequencies.

Feature Vector Vector containing all the features. For an image a feature vector could be all the raw values of pixels as a single sequence. For a trigram model with 300 dimensional word embeddings a feature vector would be a 900 dimensional vector formed by concatenating all the separate word embedding vectors.

Gazzetteer In Named Entity Recognition a gazzetteer is a dictionary of known named entities.

Language Model Probability distribution over sequences of words. Given such a sequence, say of length m, it assigns a probability $P(w_1, \ldots, w_m)$ to the whole sequence. Problems caused by growing vocabulary can be addressed with continuous language models such as neural net language models (NNML). Word2Vec by Mikolov, Corrado, et al. 2013 addresses this problem with Continuous Bag-of-words and Skip-gram models.

Lemmatisation Process of finding the base form of a word, e.g. flew -> fly

Lexeme A basic lexical unit of a language consisting of one word or several words, the elements of which do not separately convey the meaning of the whole.

n-gram Probabilistic language model where probability of current word is the joint probability of previous n words. Bigram example: $P(I, saw, the, red, house) \approx P(I|{\wedge})P(saw|I)P(red|the)P(house|red)P(\$|house)$. The words unigram, bigram and trigram language model denote n-gram model language models with n = 1, n = 2 and n = 3, respectively.

One-hot Group of bits which the legal comibations of values are only those with a single high (1) and all the others low (0). See also 1-of-V Coding.

Parsing Within computational linguistics the term is used to refer to the formal analysis by a computer of a sentence or other string of words into its constituents, resulting in a parse tree showing their syntactic relation to each other, which may also contain semantic and other information.

PMI Word co-occurence probability metric. High values for words that occur often together.

POS-tagging Process of marking up a word to particular part-of-speech (nouns, verbs, etc...) based on both its definition and its context.

ReLU Rectified Linear Unit. $h(x) = max\{0, x\}$. Used as non-linear activation function in neural nets particularly in convolutional net

Skip-gram Language model which predicts the context (previous and next n words) of a current word from the current word instead of traditional way of predicting current word from the context.

Structured Prediction Predicting structured objects, rather than scalar discrete or real values. Translating a natural language sentence into a syntactic representation such as a parse tree can be seen as a structured prediction problem in which the structured output domain is the set of all possible parse trees.

Token A structure representing a lexeme that explicitly indicates its categorization for the purpose of parsing. In plain words tokens are instances of words in a text. Not to be confused with word type.

Tree bank Parsed text corpus that annotates syntactic or semantic sentence structure. Contains trees for sentences where phrases in a sentence are structured in a tree of syntactic or semantic relations. Very useful for training POS-taggers etc...

Tri-Training Parsing unlabeled data with two different parses and selecting only the sentences for which the two parsers produce the same trees Weiss et al. 2015

Word Lookup Table Matrix $\mathbf{P} \in \mathbb{R}^{d \times |V|}$ of d rows and |V| columns, where d is the word vector dimensionality and |V| is the size of vocalbulary. I.e. each column represent a single word and each row represent single dimension in vector space.

Word Type Unique words in a text. *Good wine is good* has 4 tokens but only 3 word types.

Word vector N-dimensional vector representation of a word with interesting properties such as: vector('Paris') - vector('France') + Vector('Italy') -> vector('Rome')

0.3 Universal Dependecies

0.3.1 CoNNL-U format

Universal dependencies use CoNNL-U format for treebanks, CoNNL-U is revised version of CoNNL-X. Annotations are encoded in text files with word lines, blank lines for sentence boundaries and comments starting with hash (#).

Word lines consist of following columns:

ID	Word ID in sentence
FORM	Word form or punctuation symbol
LEMMA	Lemma or stem of word form
UPOSTAG	Universal part-of-speech tag
XPOSTAG	Language specific part-of-speech tag
FEATS	List of morphological features
HEAD	Head of the curren token, value of ID or zero (0)
DEPREL	Universal dependecy relation to the HEAD
DEPS	List of secondary dependencies
MISC	Any other annotation

Example in Finnish: Jäällä kävely avaa aina hauskoja ja erikoisia näkökulmia kaupunkiin

ID	FORM	LEMMA	UPOSTAG	XPOSTAG
1	Jäällä	jää	NOUN	N
2	kävely	kävely	NOUN	N
3	avaa	avata	VERB	V
4	aina	aina	ADV	Adv
5	hauskoja	hauska	ADJ	A
6	ja	ja	CONJ	C
7	erikoisia	erikoinen	ADJ	A
8	näkökulmia	näkö#kulma	NOUN	N
9	kaupunkiin	kaupunki	NOUN	N
10	.	.	PUNCT	Punct

FEATS
Case=Ade|Number=Sing
Case=Nom|Number=Sing
Mood=Ind|Number=Sing|Person=3|Tense=Pres|VerbForm=Fin|Voice=Act
—
Case=Par|Degree=Pos|Number=Plur
—
Case=Par|Degree=Pos|Number=Plur
Case=Par|Number=Plur
Case=Ill|Number=Sing
—

HEAD	DEPREL	DEPS	MISC
2	nmod	_	_
3	nsubj	_	_
0	root	_	_
3	advmod	_	_
8	amod	_	_
5	cc	_	_
5	conj	8:amod	_
3	dobj	_	_
8	nmod	_	SpaceAfter=No

0.3.2 Universal POS tags

ADJ	Adjective. Describing word qualifying noun or noun phrase. **deep**, **intelligent**
ADP	Adposition. Word expressing spatial or temporal relations **under**, **around**, **before** or mark various semantic roles **of**, **for**
ADV	Adverb. Modifies another word. Typically express manner, place, time, frequency etc. She sang **loudly**. You are **quite** right.
AUX	Auxiliary verb. A verb used in forming the tenses, moods, and voices of other verbs. **Do** you want tea?. He **has** given his all.
CONJ	Coordinating conjunction. Conjunction placed between words, phrases, clauses or sentences of equal rank. **and**, **but**, **or**.
DET	Determiner. Expresses reference of a noun (group). **The** girl is **a** student. **Which** book is that?
INTJ	Interjection. Shows emotion or feeling of the author, includes exclamations, curses, greetings and such. **Ouch!**, **hey**, **huh?**.
NOUN	Noun. Denotes a person, animal, place thing or idea. The **cat** sat on a **mat**.
NUM	Numeral. Number, written with digits or letters. **12**, **eleven**.
PART	Particle. Cannot be inflected. Interjections and conjunctions. In finnish also **että**, **jotta**, **koska**, **kun** etc...
PRON	Pronoun. Replaces (often previously introduced) noun. Joe saw Jill, and **he** waved at **her**.
PUNCT	Punctuation. Full stop, comma, bracket etc.
SCONJ	Subordinating conjunction. A conjunction that introduces a subordinating clause, e.g. **although**, **because**, **whenever**.
SYM	Symbol.
VERB	Verb. Conveys an action **bring**, **read**, an occurrence **happen**, **become**, or a state of being **be**, **exist**.
X	Other

0.4   Machine Translation

- See section 2 of Kestemont et al. 2016 for sources on Internet to standard language translitteration

- NLP tools suffer with texts with a lot of ortographical variation, one solution is to translate them. Kestemont et al. 2016

0.5   Language Modeling

- A language model is a probability distribution over a sequence of words, traditionally performed with n-th order Markov assumption or n-gram counting and smoothing (Chen and Goodman, 1998) *see Kim et al. 2016 Introduction for citation.*

- Kim et al. 2016 have character level encoding and word-level predicting model for language modeling.

- Kim et al. 2016 noticed that character inputs are sufficient for language modeling.

- Character level model of Kim et al. 2016 outperfom word-level and morpheme-level models on morphologically rich languages Arabic, Czech, French, German, Spanish and Russian.

- Neural Language Models NLM are blind to sub-word level information (e.g. morphemes). Kim et al. 2016

-

0.6   Intent Recognition and Slot Detection

- Intent prediction is determining user's intents from their utterances (messages). Intents can be seen as functions to call in traditional programming.

- Slot detection is identifying relevant actionable pieces of utterance Bhargava et al. 2013. Slots can be seen as function parameters.

- Bhargava et al. 2013 reduce error rates of intent prediction by 6,7% and 8.7% for transcribed data and automatically recognized data respectively when using intents from previous messages.

- Bhargava et al. 2013 find no significant difference for slot detection by using context information.

- Performance of Bhargava et al. 2013 for intent prediction is increased from 97,1% to 97.3% on transcribed data. However they use Viterbi algorithm with full access to future of the dialog, not realistic in production.

- Most systems assume a single intent per utterance leading to unnatural dialogue experience. Xu and Sarikaya 2013

- Approaching multi-intent recognition by selecting top-K hypotheses from a single intent classifier yields poor results. Xu and Sarikaya 2013

- Multi-label learning works better. Splitting into a K binary classifiers, or combining multiple labels into a single label classification problem. Xu and Sarikaya 2013

- Usually K is a system design choice. Xu and Sarikaya 2013

0.7 GloVe

- GloVe by Pennington et al. 2014 capture global corpus statistics with log-bilinear co-occurence count model.

- Memory requirements for GloVe are substantial since global co-occurence matrix for entire vocabulary is required, even though GloVe eliminates the need for zero occurence elements. Problem becomes worse for inflectional languages such as Finnish were vocabulary requires word type for each infliction for each word.

- GloVe outperforms other methods on almost all tested tasks. All tasks are English only. Pennington et al. 2014

0.8 Deep Learning

0.9 Vanishing Gradients

- Activation function with plateu will introduce gradients vanishing to zero

- RNNs have vanishing gradients problem when sequences are long.

0.10 Activation Functions

- Sigmoid has a problem with vanishing gradients when using multiple hidden layers

- Add sigmoid derivative plot for visual evidence.

- ReLU fixes this problem (still has exploding gradients problem)

- Softplus is the smooth version of ReLU but computationally more expensive.

- ELU is like ReLU but doesn't die off to zero. Clevert et al. 2015

0.11 Recurrent Neural Networks

- Multi-layer perceptrons cannot be used to map sequences to sequences since they require the dimensionality of inputs and outputs to be fixed. Sutskever et al. 2014

- RNN is suited for medling sequential phenomena. Kim et al. 2016

- RNN is a neural network architechture to which input sequence is fed one timestep at a time. RNN predicts output after each timestep and also feeds the output of previous timestep as an input in current iteration.

- In theory RNN can summarize all historical information, but in practice vanilla RNN performs poorly with long sequences due to vanishing/exploding gradients. (Bengio, Simard and Frasconi 1994), *see Model chapter of Kim et al. 2016*

- Long short-term memory (LSTM) networks address the problem of vanishing gradients with long sequences by adding a memory cell. (Hochreiter and Schmidhuber 1997), *see Model chapter of Kim et al. 2016*

- LSTM architecture addresses the vanishing gradients problem by introducing internal memory to RNN cell which is updated on every timestep via forget and update gates.

- On each timestep the LSTM forgets things from the sequence which it deems unnecessary and adds new relevant things from current example.

- Gradient exploding remains a problem, but is easily addressable in practice by simple strategies such as gradient clipping. Kim et al. 2016

- Adding more layers such that input of a layer is the hidden state of previous layer is often crucial for significant performance increase. (Pascanu et al. 2013), *see Model chapter of Kim et al. 2016*

- Deep LSTM of Sutskever et al. 2014 significantly outperformed their shallow LSTM, each layer decreasing the perplexity by nearly 10%.

- Bi-directional RNN is composed of two RNN of which the first reads the sequence in forward direction and the second reads the sequence in reverse direction, hidden states are concatenated. Chung et al. 2016

0.12 Encoder-Decoder

- Introduced by Sutskever et al. 2014 and Cho et al. 2014

- Used in machine translation Chung et al. 2016

- Dual RNN architechture where 1st RNN encodes a sequence of tokens to fixed length vector and 2nd RNN decodeds that vector representation to a target sequence of tokens. Cho et al. 2014, Bahdanau et al. 2014, Sutskever et al. 2014

- Both RNNs are jointly trained to maximize conditional probability of a target sequence given a source sequence. Cho et al. 2014, Bahdanau et al. 2014

- Encoder creates a summary of the entire input sequence. Cho et al. 2014

- Decoder samples a token at a time using input sequence summary, it's own RNN hidden state(s) and/or previously generated sample(s). Cho et al. 2014, Bahdanau et al. 2014

- See figure 1 on page 2 of Cho et al. 2014 for architecture depiction.

- Can be used to generate a target sequence based on input sequence. Can also be used to score a given pair of input-output sequences. Cho et al. 2014

- RNN Encoder-decoder captures semantic and syntactic structures of phrases. Cho et al. 2014

- Encoder-decoder needs to compress all the relevant information of the sentece in a single fixed length vector. This becomes a problem when sentence length increases, larger model is required. Bahdanau et al. 2014.

- Encoder-decoder have no explicit alignment. Liu and Lane 2016

- Input and output sequences can be of different length. **Citation?**

- Chung et al. 2016 have character level encoder-decoder sequence-to-sequence machine translation system. Using sub-word level symbols in source side, full character level only in decoder.

- Chung et al. 2016 use novel RNN (bi-scale RNN) on the target side for better handling of multiple timescales++++++

- Character level decoder relaxes the problem with computational complexity (of softmax function) with large target vocabulary. Chung et al. 2016

- Using character level only encoding and decoding removes the need to know how to do segmentation of characters into words, which is a problem in models which use character level word encodings such as C2W. Chung et al. 2016

- All inflectional forms of word result in very large vocabulary, more efficient encoding can achieved with lexeme (lemma) and morphemes, but requires a lemmatizer and morphological analyser. Chung et al. 2016

- Furthermore model may not perform well with common words if the morphological form is rare. Chung et al. 2016

- Encoder-decoder used in translation from English to French gains significant performance increase when reversing the source sentence word order. Sutskever et al. 2014

- Sutskever et al. 2014 speculate that reversing the source sentence helps by making backpropagation work better with shorter dependencies of the sentences' first words. Reversing the source sentence did not deteriorate the translation performance of later parts of the sentence, as was initially believed by Sutskever et al. 2014

0.13 Attention Mechanism

- Bahdanau et al. 2014 introduced attention mechanism in neural machine translation as a solution to deteriorating performance with long input sentences. Sutskever et al. 2014 speculate that similar improvement could have been gained with simply reversing the source sentence word order.

- System of Bahdanau et al. 2014 soft searches words in source sentence for each word in target sentence.

- System of Bahdanau et al. 2014 does not try to encode the whole sentence as a fixed size vector, but instead input sentence is encoded as sequence of vectors which are weighted at the decoding time.

- See section 3.1 of Bahdanau et al. 2014 for description of decoder with attention.

- Bahdanau et al. 2014 use beam search to approximate maximum conditional probability on the trained model.

- Attention allows encoder-decoder to learn soft alignment of input and output sequences. Liu and Lane 2016

0.14 Hyperparameter Optimization

- Random search often works well

0.15 Syntaxnet

- Andor et al. 2016

- Transition based

- Locally and globally normalized

- Backpropagation through entire net

- State-of-the-Art

1. INTRODUCTION

Testing citation Andor et al. 2016

2. NATURAL LANGUAGE PROCESSING

- Natural Language Processing is vastly wide field, this thesis discusses only on the sections of NLP relevant to the experiments.

- Subfields such as sentence segmentation and sentiment analysis are out of scope of this thesis.

- Most of the NLP work has been for english.

- Cross-linguistic annotation and parsing has been a reality only after introduction of The Universal Dependencies project and SyntaxNet.

- Similarly Finnish parsing has been unreachable until the first Finnish corpus Turku Dependency Treebank Haverinen et al. 2014 and cross-linguistic parsers.

- Traditionally NLP systems are tailored to the single problem at hand with hand engineered features suited for the problem. Recently general approach has received interest where feature engineering and task specific architectures are not needed. Collobert et al. 2011, Zhang and LeCun 2015

- See Chapter 2.1 for Finnish Language quirks in Korenius et al. 2004

2.1 Feature Engingeering in NLP

- Machine learning algorithms require words to be represented quantifiable features such as IDs or real number vectors.

- Traditional feature selection requires hand engineered features.

- Engineered features hog 95% of the computation time. Chen and Manning 2014

- Traditionally words have been represented by indices. Mikolov, Corrado, et al. 2013

- Next step was to use 1-of-V coding.

- Index representation is simple as computationally cheap, making use of huge datasets possible. Simple models with huge data outperform complex models with less data. Mikolov, Corrado, et al. 2013

- See LSA and LDA for previous systems. Neural networks significantly outperform LSA in preserving linearities. LDA doesn't scale for large datasets. Mikolov, Corrado, et al. 2013

2.2   Word Embeddings

- **see section 1.2 of Mikolov, Corrado, et al. 2013 for previous work and history of word embeddings**

- Word embeddings represent words as n-dimensional vectors. Mikolov, Corrado, et al. 2013

- LSA leverages statistical information of a corpus but performs poorly on word analogy task. Pennington et al. 2014

- Skip-gram is good for word analogies but doesn't utilize corpus statistics well since vectors are trained on local context. Pennington et al. 2014

- Word embeddings try to map words with semantic similarities close to each other. Words may have several types of similarities such as *France* and *Italy* are countries but *dogs* and *triangles* are both in plural form. Mikolov, Yih, et al. 2013

- Chen and Manning 2014 use 50 dimensional word embeddings created with Word2vec.

- Chen and Manning 2014 also use embeddings for POS tags and dependecy arcs. Only embedding POS tags has clear benefit, Chen and Manning 2014 suspect that embedding arc labels have no effect since POS tags already contain the relational information.

- Word embeddings with lookup table generalize poorly with morpohlogically rich languages such as Finnish. Takala 2016

- Morphologically rich languages benefit from breaking the word into sub-parts, RNN based character level model is not compared with Stem+ending. Takala 2016

- Word embeddings obtained through neural language models exhibit the property whereby semantically close words are close in the embedding vector space. Kim et al. 2016

- Most of the word embedding libraries work on principle of fixed vocabulary where embeddings are computed for all words in vocabulary. It's difficult to handle out-of-vocabulary words since word spelling contains only a small part of the word's semantic meaning.

- Problem for morphologically rich languages can be relaxed by lemmatizing all words because lemmas don't suffer as much from the vocabulary explosion.

2.2.1 Word2vec

- Mikolov, Corrado, et al. 2013

- Can be used with datasets of billions of words

- Has two models: Continuous bag-of-words and continuous skip-gram

- Continuous bag-of-words predicts current word from the context (surrounding words)

- Continuous skip-gram predicts context (surrounding words) from current word.

- Continuous Bag-of-Words is better for small datasets, continuous skip-gram is better for large datasets.

- CBOW is better for syntax, Skip-gram is better for semantics.

- Can be used to find semantic relationships like vector('biggest') - vector('big') + vector('small') => vector('smallest')

- State of the art (as of 2013). Since several new architectures have proposed improvements such as FastText by Facebook and GloVE by Pennington et al. 2014.

2.2.2 Charater to Word

- Ling et al. 2015

- Word embeddings can be generated from character sequences with significantly better performance for morphological languages.

- Requires only single vector for each character type. Particularly good for morphological languages where word type count may be infinite.

- Orthographical and functional (syntactic and semantic) relationships are non-trivial: *butter* and *batter* are orthographically close but functionally distant, *rich* and *affluent* are orthographically distant but functionally close. Ling et al. 2015 resort to LSTM networks for learning the relationships.

- Word lookup tables are unable to generate representations for previously unseen words, as is required for morphology. Ling et al. 2015

- C2W can generate embeddings for unseen words.

- C2W is computationally more expensive than word lookup tables, but can be eased by saving word embeddings for most frequent words since the words embedding for a character sequence (word) does not change.

- During training word embeddings change but not inside a single batch, thus it is computationally cheaper to use large batches for training.

- C2W can be replaced with word lookup tables for downstream processing since input and output of both methods are the same.

2.3 Annotations

- Stanford Dependencies by De Marneffe et al. 2006

- Stanford Dependencies emerged as de facto annotation scheme for english, but has been adapted to several other languages including Finnish. Nivre et al. 2016, Haverinen et al. 2014.

- Turku Dependency Treebank has been tranformed into universal dependencies. Pyysalo et al. 2015

- Unified annotation scheme reduces need for cross-language adaptations in downstream development. Petrov et al. 2012

- Universal Dependencies project started from the requirement for cross-linguistically consistent treebank annotations even for morphological languages. Nivre et al. 2016.

- Universal Dependencies project was born from merging several previous attempts to form a cross-linguistically sound dependency annotation schemes. Nivre et al. 2016

- UD data has been encoded in the CoNLL-U format, a revision of the popular CoNLL-X format. Nivre et al. 2016

- UD treebanks released in November 2015. Nivre et al. 2016

2.3.1 Turku Dependecy Treebank

- Haverinen et al. 2014

- Treebanks are needed in computational linguistics.

- First Finnish treebank.

- Open licence, including for text annotated

- 204339 tokens, 15126 sentences

- Based on Stanford Dependency scheme with minor modifications to exclude phenomena not present in Finnish and to include new annotations not present in English.

- Transposed to CoNNL-U scheme by universal dependencies project

- Connexor Machinese Syntax is the only currenty available Finnish full dependency parser.

- Texts from 10 different categories ranging from news and legal text to blog entries and fiction.

- Dependency parsing is done manually with full double annotation process.

- Uses Omorfi for morphological analysis. Ambiguous tokens are handled partly manually, partly rule based and partly with machine learning.

- FTB uses 3 different taggers for morphology, check them out!

- FTB is 97% grammar examples, meant for rule based POS tagger development

2.4 Tokenization

- Rule based approach to tokenization would mostly split sentence into words from spaces and separate punctuation from the words.

- Symbols and codes are more challenging for rule based tokenizers.

- Neural net based approach to tokenization can be done with seq2seq model which inserts linefeeds.

- Another neural net approach to tokenization is to do character classification where each character is classified to be first character of a word.

- Good tokenization is very important for good downstream processing results; very small errors in tokenization can lead to extremely large errors in subsequent tasks (ask sitation from Honain).

- This thesis uses gold standard tokenization of the data files and therefore tokenization is not included in the experiments.

2.5 POS-tagging

- Started from rule based taggers

- Tagger by Brill 1992 (known as Brill tagger) learns the rules and as such can be considered as a hybrid approach

- Contemporary research is focused on statistical and NN based taggers

- Rest of this section focuses on statistical parsers

- Ling et al. 2015 introduced S-LSTM based State-of-the-art tagger

- Andor et al. 2016 Improved accuracy with transition based tagger

- Chen and Manning 2014 were first to represent POS-tag and arc labels as embeddings

- Andor et al. 2016 and Weiss et al. 2015 built their solutions based on Chen and Manning 2014

- Nivre 2004 introduced system for transition based taggers known as arc-starndard system. Chen and Manning 2014

2.6 Lemmatisation

- Lemmatization is the process of finding a base form for a word.

- Lemmatization is a normalization technique. Korenius et al. 2004

- Homographic and inflectional word forms cause ambiquity. Korenius et al. 2004

- Compound words cause problems. Korenius et al. 2004

- Lemmatization is better than stemming for clustering of documents written in Finnish because of it's highly inflectional nature. Korenius et al. 2004

- Lemmatization catches better the semantic meaning of a word, as can be deducted from a better clustering performance.

- Omorfi does lemmatization based on morphological analysis

- Omorfi produces multiple lemmas which need to be disambiguated

- Disambiguation can be done with selecting most probable word, given the context, by language model

- Kestemont et al. 2016 try to solve lemmatization as a neural net classification problem, where lemmas are the class labels

- Method of Kestemont et al. 2016 cannot produce lemmas not seen on training time.

- Lemmatization has received a lot of research attention for highly inflectional languages, see Kestemont et al. 2016

- Lemmatization of english is considered a solved problem, rule based or hybrid approaches can do practically flawless job.

- There has been almost none previous work using deep learning for lemmatization before Kestemont et al. 2016

2.7 Morphological Parsing

TODO: Should this be included at all?

2.8 Structural Parsing

TODO: Should this be included at all?

- Aims to find structure of a sentence.

- Commonly divided to two different tasks: constituency parsing and dependency parsing.

- Constituency parser creates a parse tree of constituencies.

- Dependency parser creates a parse tree of word token dependencies.

- Constituency parsers are slower but more informal than dependency parsers. Fernández-González and Martins 2015

- Fernández-González and Martins 2015 show that it is possible to build constituency parser with dependency parser by reducing constituents to dependency parsing.

- CoNLL uses dependency parse trees.

2.8.1 Transition Based Parsers

- Good balance between efficiency and accuracy Weiss et al. 2015

- Parsed left to right; at each position the parses chooses action from a set of possible actions.

- Greedy models are fast but error prone and need hand engineered features Weiss et al. 2015

- Actions can be chosen by ANN to avoid hand engineering Chen and Manning 2014, Weiss et al. 2015

3. EXPERIEMENTS ON JOINT MODEL FOR POS-TAGGING AND LEMMATIZATION

Focus of this thesis is to prove or disprove the hypothesis that joint learning of lemmatization and POS-tagging with neural networks can obtain better performing network for one of both tasks than learning the said tasks separately. Joint learning in the context of this thesis means learning and predicting both outputs with single neural network architecture and single forward pass. Joint learning model is compared to separate tasks baseline.

Lemmatization and POS-tagging were selected as tasks for this thesis because they are well studied and results from other research projects exist making baseline validation possible. Both tasks are also popular choises as input features for downstream processing. Although lemmatization is considered by some to be solved for morphologically poor languages, it is not for morphologically rich languages such as Finnish. Lemmatization is especially interesting for Finnish because properly done lemmatization would allow usage of word level features such as pre-computed word2vec vectors as input features in downstream tasks. All experiments of this thesis are performed with Finnish language.

Neural networks were selected as implementational approach for the problem mainly because of their flexible and architecturally general nature. Neural networks don't require architectural changes, other than maybe a hyper-parameter optimization, when adapting the network for new languages. Essentially same neural network can handle the natural language processing task at hand for any language. Possible exceptions are languages which are written on different level than european languages. Chinese has symbols only for words, has no letters at all, and as such might demand architectural changes.

Another convenient property of neural networks is their flexibility to adapt different tasks with sometimes very small architectural changes. Neural networks developed for this thesis share vast majority of components among lemma classifier and POS classifier, only the output layer is separate and has different number of nodes for said tasks. But even then both outputlayers are similar fully connected linear projection

layers.

Lastly neural networks were selected because NLP has been researched for several decades and it appears that older, often statistical, methods have been already tuned close to their maximum. Neural networks on the other hand have shown very promising progress during the last couple of years, mainly due to ever decreasing price of computational resources and introduction or re-introduction of a few mathematical advancements which have made deep neural networks easier to train.

Statistical methods such as bag-of-words might be trickier to implement as character-level models than neural networks. Counting words in a sentence of in a context of a word has for a years been the simplest baseline model for mulitude of NLP tasks. Bag-of-words works because words are essentially the basic sematic unit of a language. Character on the other hand convey very little meaning when not associated with other characters in order of appearance. Simply counting characters is therefore not going to reveal the underlying phenomena. Same explanation applies to some extent to other traditional models also.

Joint model approach was taken into inspection because joint learning models have not yet been studied very widly but results from the few studies have shown that joint models can achieve better results than separate models. Lemma and a part of speech of a word a tighly linked to each other. Lemma is the basic (almost) unique identifier of a word and each word has always a single fixed part-of-speech. This tight coupling of lemmas and part of speeches serve as a good foundation for building a joint learning model.

Unfortunately lemmas do not identify a word in a completely unique way; multiple words may have same written base form. Nail is a written form for at least two different meanings, one being a fastener for attaching pieces of wood together with a hammer and the other being a keratin made envelope covering the tips of fingers and toes. Fortunately words with multiple meanings are more rare than not. This thesis simply omits the problems and implications which could and do arise from having shared tokens for multiple words. Such decision to omit the problems may not be as harmful as one might think; predicting correct lemma whether the nail is meant to be hit with a hammer or not does still produce a correct lemma, this becomes and issue only with the downstream tasks which might need the two to be separated.

Part of speech tagging can also still be done without too much hinderance. Often the two words have the same part of speech. In the case of fastener nail and finger nail problem does not exists since both are nouns. Separating "nails" as plural form

of a noun "nail" from "nails" as colloquial form of a verb for love making can be done with the information provided by the context of the word.

3.1 Neural Network Architechture

TODO: Architecture diagram

Architecture used for experiments in this thesis is a multi-layer deep neural network composing an end to end pipeline handling everything needed from character representations to output classifications. Multiple layers in the architecture are not only layers of interconnected nodes as usually depicted by term layers in the context of neural networks. The architecture also contains multiple architectural layers such as character embedding layer, word embedding layer, context encoding layer and classification layer. To distinguish architectural layers from neural layers the former is going to be called components for the rest of this thesis. All the components, all layers which make the individual components and weights are learned at the same time. Learning all network parameters in a single training run makes training simpler and removes, or at least obfuscates, the possible compatibility issues and non-optimalities between the components and layers.

Neural network code was implemented with Python using popular computation libraries Tensorflow and Numpy. Tensorflow provides scripting APIs for other programming languages too but Python was selected because of it is fast to write and computational performance does not suffer compared to eg. C++ since all the expensive computations are made in the C++ based backend. Numpy is used in the data pre-processing phase where data is processed to be suitable for feeding as input to the actual neural network. Tensorflow was used for implementing the neural network as a computational graph. Tensorflow makes it fairly easy to implement compled multi-componen end to end architectures while abstracting the actual optimization work away from the developer.

Tensorflow also provides a layer of abstraction to execute the computational graph on either CPU or GPU. GPU computations for pleasantly parallel problems such as neural network optimization are significantly faster. Doing the network training and especially hyper-arameter optimization only on CPU would have not been feasible with the hardware resources available. If computations would have been forced to be ran on CPU, there would have not been enough resources to run sufficient number of hyper-parameter optimization runs. Neural network results are often very sensitive to having suitable hyper-parameters and therefore results obtained in the experiments in the worst case could have been unable to provide answer for the

hypothesis.

3.1.1 Lemmatizations as Classification Only Task

Approach for lemmatization in this thesis is classification only. POS-tags are a limited set of 31 labels as defined in the Universal Dependencies project but such is not the case for lemmas. When classifying lemmas with neural network one needs an indexed vocabulary of all possible output labels. Having fixed and limited set of lemmas proves to be a challenge for lemmatization.

As discussed earlier the Finnish language suffers from vocabulary explosion even when considering only lemmas and omitting the inflections because Finnish language makes it possible to form compound words very freely. The number of possible combinations created by selecting two or more words for a compound word is way too large to be handled with a linear vocabulary. The word2vec vocabulary which is created from Finnish Internet Parse bank TODO: citation contains over 1,7 million unique lemmas. Learning to classify this number of lemmas with a corpus of 160 and some thousand tokens is obviously impossible task.

To circumvent the vocabulary explosion problem for output vocabulary fixed and limited set of lemmas were selected from the training set. Selected lemma vocabulary contains 90% of the use cases in the Finnish Universal Dependencies 1.4 training dataset. The 90% coverage is formed by selecting the most frequent words only. This lemma vocabulary contains less than nine thousand unique lemmas opposed to over 1,7 million in the word2vec vocabulary. All the lemmas that were left out of the selected vocabulary are treated as unknown tokens meaning that when classifying out of vocabulary lemma the neural network with output an unknown token "<UNK>". Since vocabulary covers 90% of the uses in training set, 10% of uses are left out making unknown token a most frequent label.

Having unknown tokens provides it's own set of challenges for downstream processing tasks: information that is supposed to be provided by the lemma is lost with unknown token. One approach for outputting all possible lemmas is to use a generative model such as encoder-decoder model popular in recent studies TODO: citation for machine translation. Generative model does not classify indices to a fixed vocabulary but generates the lemma one character at a time.

Generative model does solve the problem with lemma vocabulary but introduces a myriad of other problems. Some of the most prominent problems being vastly increased architectural complexity and significantly increased computational complex-

ity and memory requriements. Introducing a encoder-decoder model also introduces problems with observation metrics. POS-tagging and lemma classification are word level prediction tasks and are as such measured with word level metrics eg. accuracy or F1 score. However encoder-decoder model is a character level model which is also observed on character level. Mixing word level classifications and character level classifications fuzzies the meaning of used metrics for evaluation of network performance.

Because of the forementioned problems, the encoder-decoder model was not a part of the implementation used for this thesis and lemmatization is treated as a classification only task. Also it's worth noting that observing lemma classification with POS-tagging should prove to be sufficient for proving or disproving the hypothesis.

3.1.2 Word Embedding

- C2W

- Character vocabulary needs to fixed for character embedding layer.

- We selected `!"#$%&'()*+,-./0123456789:;<=>?@ABCDEFGHIJKLMNOPQRSTUVWXYZ[\]^_`abcdefghijklmnopqrstuvwxyz{|}~ÄÅÖäåö€` as our character vocabulary.

- This contains ASCII characters from 32 to 127, scandic characters in lower and upper case and € sign.

- Selected character vocabulary covers 99,933% of character usages in Finnish internet parsebank n-gram dataset.

- ASCII end-of-text (ETX) and substition (SUB) characters were added to character vocabulary to handle sequence length padding and out of vocabulary characters respectively.

- Characters were represented as character embedding vectors obtained during the training and implemented as simple trainable Tensor in tensorflow.

- C2W RNN takes a mini-batch of words represented as character tensors (one character per row) as an input

- Output of C2W is word embedding vectors for all words in the mini-batch.

3.1.3   Context Encoding

- Bi-directional RNN

- Word vectors obtained from word embedding layer used as input

- Input for each word also contains it's right and left side contexes.

- Each word has effectively only uni-directional RNN for each side context because sentence is fed once through bi-directional RNN. Forward pass encodes left side contexes and backward pass encodes right side contexes.

- Uni-directional RNN encoding for contexes should not be a problem; sentence lengths are usually very limited (averaging 13 words) and most important words in the context are the words closest to current word for whom context is being encoded.

- Output is word embedding vector which also contains information about the word's context.

3.1.4   Classification

- Add fully connected layers and output layer (MLP)

- Decision to add fully connected layers was made with brief experiments comparing results with and without the fully connected layers. Fully connected layers added almost no computational complexity while improving learning.

- All experiments were done with fully connected layers and further experimenting how they affect different tasks was left for future research.

- Lemma vocabulary consists of most frequent lemmas in Universal Dependencies dataset covering at least 90% of use cases in the dataset.

- Output layers are a simple fully connected layers with linear activation functions. Size of output lemma classification and POS classification layers are the size of lemma vocabulary and size of POS-tag vocabulary respectively.

3.1.5   Training and Optimization

- Back-propagation by Tensorflow's innate graph functionality

- Weights optimized with AdamOptimizer (gradient descent)

- Hyperparameters optimized with Optunity

- Hyperparameter optimization algorithm: Particle Swarm

- Time budget of one hour per hyperparameter optimization run

- Early stopping on each optimization run based on validation loss

- Fixed number of optimization runs

- Trained on NVIDIA GeForce TitanX (Pascal)

3.2 Experiments

- Minimal set for testing the hypothesis

- POS-tagging without lemmatization

- Lemmatization without POS-tagging

- Lemmatization with classified POS-tags

- Lemmatization with gold standard POS-tags

- Joint model

- Comparing these results should be enough to (dis)prove the hypothesis

- Not testing on other languages

- Not testing how tokenization affects results

- UDPipe tokenization accuracy is 99,7% Straka and Straková 2017

- POS-taggin with lemmas not tested

- POS-tagging with lemmas as input features is not very meaningful task because POS-tag can be read from a dictionary if lemma is known.

- Not testing with other vocabulary coverages

- Not testing with XPOSTAGS

- Not testing morphology

3.2.1 Test Methods

TODO: Excerpt from Finnish UD data

- Testing with Universal Dependencies data

- Finnish UD data is tagged with Omorfi and manually corrected, see 2.3.1 TODO: Check manual correction

- UD has pre-existing results from third parties such as Syntaxnet and UDPipe for baseline comparisons

- Other annotated corpuses exist for Finnish but UD is the only multi-lingual corpus collection where all datasets are in uniform format making future research and comparison easier. TODO: citations!

- Using accuracy. Percentage of correct classification labels.

- Academia does not seem to agree what F1 score is

- Accuracy is simple and easier for the readers to understand

- Accuracy may not be suitable if neural network prediction confidences are important. Neural net may learn to be very confident about wrong predictions if only accuracy is observed.

- Accuracy is suitable to distinguish between results on different configurations. Confidences are not needed for this work.

- Cross-entropy based loss is good for weight optimization and does not suffer from confidence issues but tells very little to a reader without deep understanding about the task at hand.

- Each experiement was executed by training the network on the training dataset with early stopping based on observations of validation dataset loss.

- Same hyperparameters were used for all experiements, tuned with Optunity implemented particle swarm algorithm for joint model.

- All results reported are tested with test dataset.

- Tokenized data (forms) is fed to neural network as input and neural network predictions are compared to correct values.

3.3 Results

- UDPipe 1.2 achieves 94,9% on POS-tagging on Finnish Universal Dependencies 2.0 and 86,8% on lemmatization. Both results with gold tokenization. Straka 2017

- UDPipe lemmatizer is a generative model (Straka and Straková 2017), so comparison is not fair or even meaningful.

- POS-tagging without lemmatization = 94,30%. Remarkably close to UDPipe results with only 0,6%-points advantage to UDPipe. Limited mainly by data?

- Lemmatization without POS-tagging = 94,25%

- Lemmatization with classified POS-tags as input features = 94,40%

- Lemmatization with gold standard POS-tags as input features = 96,22%

- Joint model. Lemmas = 95,24%, POS = 94,14%

- Joint model achieves best practical results for lemmatization, with 0,99%-points increase, since gold-standard POS-tags are not available for live inferencing.

- Joint model POS-tags are slightly inferior (0,16%-points) to baseline

- Lemmatization results with classified POS-tags are slightly better (0,15%-points) than the baseline.

- Lemmatization with gold-standard POS-tags achieves clearly the best results with 1,97%-points increase over baseline.

4. DISCUSSION

4.1 How well results generalize?

- Do these results generalize to other languages needs to be verified with additional experiments.

- Universal Dependencies datasets are created to provide as uniform language modeling as possible across all languages and therefore it is probable that the same mutual information between lemmas and POS-tags in other languages would prove beneficial in joint learning of lemma and POS-tag classification.

- Speculating how well these results generalize to other tasks in natural language processing is more complicated because other tasks may not share as strong mutual information as lemmas and POS-tags do.

- Liang et al. 2016 proved that neural networks benefit from jointly learning two tasks different from lemmatization and POS-tagging. Liu and Lane 2016 even had tasks on different architectural level, slot filling being word level task and intent classification being sentence or message level task.

- With results provided in this work and results from Liu and Lane 2016 it's failry safe to say that neural networks can obtain better results when learning jointly two tasks that share mutual information.

4.2 What was assumed?

- It was assumed that training dataset would represent Finnish language well. May not hold true since all sources of the dataset are from internet.

- It was also assumed that different datasets would have fairly similar data distribution and training done on training dataset would generalize to test dataset.

- It was noticed that selecting lemma vocabulary which covers 100% of uses in training set only covers about 75% of uses in validation and test datasets.

- Lemma classifier learned remarkably well to classifiy lemmas not seen during training time as unknown tokens. Maybe priori bias explains this?

4.3 What was simplified?

- Slightly different results would be optained by optimizing hyperparameters separately for every different experiement. Hyperparameter optimization requires about 100 to 200 training runs and as such takes significant time budget. Using time required for hyperparameter optimization before every experiment proved inconvenient for this work.

- Hyperparamenter optimization do not have variation with loss weights for the two tasks. Loss weight sets the target if optimizing based on validation loss and therefore cannot be optimized.

- Hyperparameter optimization was using time budget for each run

- Hyperparameter optimization had limited number of runs available

- Selecting a different optimization target for hyperparameters could have allowed loss weight optimization.

- Character vocabualry for C2W does not contain all characters in datasets. Even still embedding vectors for some characters are probably not well learned because of their unfrequency.

- Lemmatization as word level classification task only is also a pretty significant simplification. Decoder model for lemma generation is able to always produce some kind of answer without ever producing unknown tokens.

- Selected lemma vocabulary covers 90% of the training set -> unknowns are 10%, second most frequent is less than 5%. It was not tested if unknown prediction accuracy would break if unknowns weren't the most frequent.

- Whether having unknown tokens is fatal flaw depends on the problem at hand and on the down stream processing which uses lemmas obtained with lemmatization. For some tasks not having generated lemmas with high probability of typographical errors produced might be significantly better option.

- Experiements done with golds standard tokenization of Universal Dependencies datasets. In live production such tokenization is not available. Imperfect tokenization may completely ruin the classification if tokens are not created correctly.

- Same applies for sentence segmentation.

5. CONCLUSIONS

- Hypothesis proved to be true, at least partly

- Neural networks are very suitable for joint learning tasks.

- Uncorrelated tasks are not tested yet but those don't have any obvious theoretical foundation why they should benefit.

- Interestingly we couldn't obtain increase in POS-tagging performance.

- Maybe lemmatization as a more information rich task dominates the learning

- POS-tagging on joint model could be improved by giving more weight to POS-tagging task but that quickly deteriorates the lemmatization perfomance drastically.

- Doing neural net based lemmatization benefits from jointly learning also to POS-tag. Added architectural, implementational and computational complexities are minor.

- Author recommends always to do POS-tagging when doing lemmatization.

- If main task is POS-tagging, adding lemmatization is probably not beneficial enough for the added computational cost.

- If both tasks need to be done and POS-tagging performance is a priority it might prove to be difficult to achieve best possible results with joint model.

BIBLIOGRAPHY

Andor, D. et al. (2016). "Globally Normalized Transition-Based Neural Networks". In: *Acl 2016*, pp. 2442–2452. DOI: `10.18653/v1/P16-1231`. arXiv: `arXiv:1603.06042v2`.

Bahdanau, D., K. Cho, and Y. Bengio (2014). "Neural Machine Translation By Jointly Learning To Align and Translate". In: *Iclr 2015*, pp. 1–15. ISSN: 0147-006X. DOI: `10.1146/annurev.neuro.26.041002.131047`. arXiv: `1409.0473`. URL: `http://arxiv.org/abs/1409.0473v3`.

Bhargava, A. et al. (2013). "Easy contextual intent prediction and slot detection". In: *ICASSP, IEEE International Conference on Acoustics, Speech and Signal Processing - Proceedings*, pp. 8337–8341. ISSN: 15206149. DOI: `10.1109/ICASSP.2013.6639291`.

Brill, E. (1992). "A Simple Rule-Based Part of Speech Tagger". In: *Applied natural language*, p. 3. ISSN: 00992399. DOI: `10.3115/1075527.1075553`. arXiv: `9406010 [cmp-lg]`.

Chen, D. and C. D. Manning (2014). "A Fast and Accurate Dependency Parser using Neural Networks". In: *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)* i, pp. 740–750. ISSN: 9781937284961. URL: `https://cs.stanford.edu/%7B~%7Ddanqi/papers/emnlp2014.pdf`.

Cho, K. et al. (2014). "Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translation". In: *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 1724–1734. ISSN: 09205691. DOI: `10.3115/v1/D14-1179`. arXiv: `1406.1078`. URL: `http://arxiv.org/abs/1406.1078`.

Chung, J., K. Cho, and Y. Bengio (2016). "A Character-level Decoder without Explicit Segmentation for Neural Machine Translation". In: *Acl-2016*, pp. 1693–1703. arXiv: `1603.06147`.

Clevert, D.-A., T. Unterthiner, and S. Hochreiter (2015). "Fast and Accurate Deep Network Learning by Exponential Linear Units (ELUs)". In: *Under review of ICLR2016 ELU* 1997, pp. 1–13. arXiv: `1511.07289`. URL: `http://arxiv.org/pdf/1511.07289.pdf%7B%5C%%7D5Cnhttp://arxiv.org/abs/1511.07289`.

Collobert, R. et al. (2011). "Natural Language Processing (Almost) from Scratch". In: *Journal of Machine Learning Research* 12, pp. 2493–2537. ISSN: 0891-2017. DOI: `10.1.1.231.4614`. arXiv: `1103.0398`.

De Marneffe, M.-C., B. MacCartney, and C. D. Manning (2006). "Generating typed dependency parses from phrase structure parses". In: *Proceedings of the 5th International Conference on Language Resources and Evaluation (LREC 2006)*,

pp. 449–454. DOI: 10.1.1.74.3875. URL: http://nlp.stanford.edu/pubs/LREC06%7B%5C_%7Ddependencies.pdf.

Fernández-González, D. and A. F. T. Martins (2015). "Parsing as Reduction". In: *In Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)* 1983, pp. 1523–1533. arXiv: 1503.00030.

Harris, Z. S. (1954). "Distributional Structure". In: *WORD* 10.2-3, pp. 146–162. DOI: 10.1080/00437956.1954.11659520. URL: http://dx.doi.org/10.1080/00437956.1954.11659520.

Haverinen, K. et al. (2014). "Building the essential resources for Finnish: the Turku Dependency Treebank". In: *Language Resources and Evaluation* 48.3, pp. 493–531. ISSN: 15728412. DOI: 10.1007/s10579-013-9244-1.

Kestemont, M. et al. (2016). "Lemmatization for variation-rich languages using deep learning". In: *Digital Scholarship in the Humanities*, fqw034. ISSN: 2055-7671. DOI: 10.1093/llc/fqw034. URL: http://dsh.oxfordjournals.org/lookup/doi/10.1093/llc/fqw034.

Kim, Y. et al. (2016). "Character-Aware Neural Language Models". In: *Aaai*. arXiv: 1508.06615. URL: http://arxiv.org/abs/1508.06615.

Korenius, T. et al. (2004). "Stemming and lemmatization in the clustering of finnish text documents". In: *Proceedings of the thirteenth ACM conference on information and knowledge management*, pp. 625–633. DOI: 10.1145/1031171.1031285. URL: http://portal.acm.org/citation.cfm?id=1031171.1031285%7B%5C&%7Dcoll=Portal%7B%5C&%7Ddl=ACM%7B%5C&%7DCFID=88534260%7B%5C&%7DCFTOKEN=49348956.

Liang, C. et al. (2016). "Neural Symbolic Machines: Learning Semantic Parsers on Freebase with Weak Supervision". In: October. arXiv: 1611.00020. URL: http://arxiv.org/abs/1611.00020.

Ling, W. et al. (2015). "Finding Function in Form: Compositional Character Models for Open Vocabulary Word Representation". In: *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing* September, pp. 1520–1530. DOI: 10.18653/v1/D15-1176. arXiv: 1508.02096. URL: http://dx.doi.org/10.18653/v1/d15-1176%7B%5C$%7D%7B%5C%%7D5C%7B%5C$%7Dnfile:///Files/68/6810072d-e133-426e-807f-445df2840420.pdf%7B%5C$%7D%7B%5C%%7D5C%7B%5C$%7Dnpapers3://publication/doi/10.18653/v1/d15-1176%7B%5C$%7D%7B%5C%%7D5C%7B%5C$%7Dnhttp://arxiv.org/abs/1508.02096.

Liu, B. and I. Lane (2016). "Attention-Based Recurrent Neural Network Models for Joint Intent Detection and Slot Filling". In: 1, pp. 2–6. DOI: 10.21437/Interspeech.2016-1352. arXiv: 1609.01454. URL: http://arxiv.org/abs/1609.01454.

Mikolov, T., G. Corrado, et al. (2013). "Efficient Estimation of Word Representations in Vector Space". In: *Proceedings of the International Conference on Learning Representations (ICLR 2013)*, pp. 1–12. ISSN: 15324435. DOI: 10.1162/153244303322533223. arXiv: arXiv:1301.3781v3. URL: http://arxiv.org/pdf/1301.3781v3.pdf.

Mikolov, T., W.-t. Yih, and G. Zweig (2013). "Linguistic regularities in continuous space word representations". In: *Proceedings of NAACL-HLT* June, pp. 746–751. URL: http://scholar.google.com/scholar?hl=en%7B%5C&%7DbtnG=Search%7B%5C&%7Dq=intitle:Linguistic+Regularities+in+Continuous+Space+Word+Representations%7B%5C#%7D0%7B%5C%%7D5Cnhttps://www.aclweb.org/anthology/N/N13/N13-1090.pdf.

Nivre, J. (2004). "Incrementality in deterministic dependency parsing". In: *Proceedings of the Workshop on Incremental Parsing: Bringing Engineering and Cognition Together*, pp. 50–57. DOI: 10.3115/1613148.1613156. URL: http://dl.acm.org/citation.cfm?id=1613156.

Nivre, J. et al. (2016). "Universal Dependencies v1: A Multilingual Treebank Collection". In: *Proceedings of the 10th International Conference on Language Resources and Evaluation (LREC 2016)*, pp. 1659–1666.

Pennington, J., R. Socher, and C. D. Manning (2014). "GloVe: Global Vectors for Word Representation". In: *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing*, pp. 1532–1543. ISSN: 10495258. DOI: 10.3115/v1/D14-1162. arXiv: 1504.06654.

Petrov, S., D. Das, and R. Mcdonald (2012). "A Universal Part-of-Speech Tagset". In: arXiv: 1104.2086.

Pyysalo, S. et al. (2015). "Universal Dependencies for Finnish". In: *Nordic Conference of Computational Linguistics NODALIDA 2015* Nodalida, p. 163.

Straka, M. (2017). *UDPipe User's Manual*. URL: http://ufal.mff.cuni.cz/udpipe/users-manual (visited on 10/01/2017).

Straka, M. and J. Straková (2017). "Tokenizing, POS Tagging, Lemmatizing and Parsing UD 2.0 with UDPipe". In: *Proceedings of the CoNLL 2017 Shared Task: Multilingual Parsing from Raw Text to Universal Dependencies* 2, pp. 88–99. URL: http://www.aclweb.org/anthology/K17-3009.

Sutskever, I., O. Vinyals, and Q. V. Le (2014). "Sequence to sequence learning with neural networks". In: *Nips*, pp. 1–9. ISSN: 09205691. DOI: 10.1007/s10107-014-0839-0. arXiv: 1409.3215. URL: http://papers.nips.cc/paper/5346-sequence-to-sequence-learning-with-neural.

Takala, P. (2016). "Word Embeddings for Morphologically Rich Languages". In: April, pp. 27–29.

Weiss, D. et al. (2015). "Structured Training for Neural Network Transition-Based Parsing". In: *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)* 2012, pp. 323–333. DOI: 10.3115/v1/P15-1032. arXiv: 1506.06158. URL: http://www.aclweb.org/anthology/P15-1032.

Xu, P. and R. Sarikaya (2013). "Exploiting shared information for multi-intent natural language sentence classification". In: *Proceedings of the Annual Conference of the International Speech Communication Association, INTERSPEECH* 1, pp. 3785–3789. ISSN: 19909772.

Zhang, X. and Y. LeCun (2015). "Text Understanding from Scratch". In: *APL Materials* 3.1, p. 011102. ISSN: 2166-532X. DOI: 10.1063/1.4906785. arXiv: 1502.01710. URL: http://arxiv.org/abs/1502.01710.

APPENDIX A. SOMETHING EXTRA

Appendices are purely optional. All appendices must be referred to in the body text

APPENDIX B. SOMETHING COMPLETELY DIFFERENT

You can append to your thesis, for example, lengthy mathematical derivations, an important algorithm in a programming language, input and output listings, an extract of a standard relating to your thesis, a user manual, empirical knowledge produced while preparing the thesis, the results of a survey, lists, pictures, drawings, maps, complex charts (conceptual schema, circuit diagrams, structure charts) and so on.